

# SHL Assessment Recommendation System

[By KHUSHI SHUKLA]

## Approach and Optimization Document

---

### 1. Objective

The goal of this project was to design a recommendation system that maps job descriptions or user queries to the most relevant **SHL assessment products**. The system needed to support **an API**, a **Streamlit web interface**, and produce a **CSV output** of predictions on the provided test dataset.

---

### 2. Overall Approach

The solution integrates **semantic embedding models** to capture the meaning of both job descriptions and assessment content.

Key steps followed:

#### Step 1 – Data Preparation

- The file **products.csv** was created from the source Excel file **Gen\_AI Dataset.xlsx**. The original dataset contained detailed information about SHL products, including names, descriptions, and categories. To prepare it for the recommendation model, only the **text-based columns** — mainly *Product Name* and *Description* — were extracted. The text was cleaned by removing missing values, extra spaces, and special symbols.
- The provided product catalog (products.csv) was preprocessed by combining fields like *AssessmentName*, *Description*, and *Category* into a unified text representation.

#### Step 2 – Model Selection

- The primary model used was **SentenceTransformer – all-MiniLM-L6-v2**, a compact yet high-performing model optimized for semantic similarity.
- Additionally, a **Hugging Face Transformer model** was integrated to experiment with custom embeddings and compare against SentenceTransformer results.
- Gemini LLM was initially considered for cloud-based embedding generation, but due to quota limits, the final version focused on **offline models**, ensuring reproducibility and cost-free usage.

#### Step 3 – Embedding Generation

- Both the catalog entries and user/job queries were converted into **vector embeddings** using the transformer model.

- **Cosine similarity** was used to measure semantic closeness between a query and each assessment.

#### Step 4 – Recommendation Logic

- For each query, similarity scores were sorted in descending order.
  - The top-N recommendations were returned (N configurable in Streamlit).
  - The system output included the **most relevant assessment name**, its **similarity score**, and a **ranked list** of top matches.
- 

### 3. System Components

Component	Description
<b>1. Streamlit App</b>	User interface to input job descriptions, select model type, and view recommendations interactively.
<b>2. FastAPI Backend</b>	Serves as an API endpoint for programmatic access to the recommender, returning JSON responses.
<b>3. Prediction Script</b>	Generates predictions for all queries in the given test set and saves them as predictions.csv.

---

### 4. Optimization Strategy

During development, several experiments were conducted to improve the **relevance and stability** of recommendations:

Stage	Optimization	Outcome
<b>Baseline</b>	Used default SentenceTransformer embeddings with raw text.	Moderate accuracy, noisy recommendations.
<b>Text Enrichment</b>	Concatenated multiple catalog columns into combined_text.	Improved context representation.
<b>Normalization</b>	Normalized embeddings and removed stopwords.	Enhanced semantic distance accuracy.
<b>Model Fine-Tuning</b>	Adjusted sentence length truncation and batch size.	Faster inference with minimal accuracy drop.

Stage	Optimization	Outcome
Top-K Tuning	Experimented with top-3, top-5, and top-10 results.	Top-10 proved most consistent for HR domain relevance.

The final model achieved a **balanced trade-off** between speed, interpretability, and recommendation precision.

## 6. Key Takeaways

- **Offline embeddings** via SentenceTransformer are fast and reliable for semantic similarity tasks.
- **Hugging Face models** offer flexibility to extend to domain-specific fine-tuning if more labeled data becomes available.
- **Configurable recommendations (Top-K)** in Streamlit improve usability for HR professionals exploring multiple assessments.
- The pipeline is **modular**, enabling future integration with APIs, caching, or database indexing for large catalogs.

## 7. Final Deliverables

Deliverable	Description
API Endpoint	FastAPI service returning top-N recommendations in JSON format.
Web Interface	Streamlit app for interactive exploration.
CSV File	predictions.csv – Predictions for test queries with ranked top-10 results.
GitHub Repository	Complete implementation with code, model references, and documentation.

## 8. Future Improvements

- Fine-tune the transformer model using labeled SHL historical recommendation data.
- Implement **vector database indexing** (FAISS / Pinecone) for faster retrieval on larger datasets.
- Integrate a **feedback loop** to continuously improve relevance using user click data.

---

## Conclusion

This project demonstrates a robust and interpretable **semantic recommendation system** tailored for SHL assessments.

Through model selection, embedding optimization, and modular deployment (Streamlit + FastAPI), the system achieves a high degree of accuracy, scalability, and end-user usability.